

Contents

1 Chapter 0 — Motivation: Your First RL Experiment	1
1.1 0.0 Who Should Read This?	1
1.2 0.1 The Friday Deadline	2
1.3 0.2 The Core Insight: Boosts as Actions	2
1.4 0.3 A Tiny World: Toy Simulator and Reward	3
1.4.1 0.3.1 User Types	3
1.4.2 0.3.2 Products (Sketch)	3
1.4.3 0.3.3 Actions: Boost Weight Templates	4
1.4.4 0.3.4 Toy Reward Function	4
1.5 0.4 Your First RL Agent: Tabular Q-Learning	5
1.5.1 0.4.1 Problem Recap	5
1.5.2 0.4.2 Algorithm: ϵ -Greedy Q-Learning	5
1.5.3 0.4.3 Minimal Implementation	5
1.5.4 0.4.4 Learning Curves and Baselines	8
1.6 0.5 Reading the Experiment: What We Learned	9
1.6.1 Convergence Pattern	9
1.6.2 Per-Segment Performance	10
1.6.3 What We (Hand-Wavily) Assumed	11
1.6.4 Theory-Practice Gap: ϵ -Greedy Exploration	11
1.7 0.6 Limitations: Why We Need the Rest of the Book	12
1.7.1 1. Discrete Action Space	12
1.7.2 2. Tabular State Representation	12
1.7.3 3. No Constraints	12
1.7.4 4. Simplified Position Bias	13
1.7.5 5. Online Exploration Risk	13
1.7.6 6. Single-Episode Horizon	13
1.8 0.7 Map to the Book	13
1.9 0.8 How to Use This Book	14
1.9.1 For Practitioners	14
1.9.2 For Researchers / Mathematically Inclined	14
1.9.3 For Everyone	14
1.10 Exercises (Chapter 0)	14
1.11 0.9 Code Artifacts	16
1.12 0.10 What's Next?	17

1 Chapter 0 — Motivation: Your First RL Experiment

Vlad Prytula

1.1 0.0 Who Should Read This?

This is an **optional but highly recommended warm-up**. It’s deliberately light on mathematics and heavy on code. You’ll build a tiny search world, train a small agent to learn context-adaptive boost weights, and see the core RL loop in action. The rest of the book (Chapters 1–3) provides the rigorous foundations that explain *why* this works and *when* it fails.

If you're a practitioner: Start here. Get your hands dirty with a working system in 30 minutes. Return to the theory when you need deeper understanding.

If you're mathematically inclined: Skim this chapter to see where we're headed, then dive into Chapter 1's rigorous treatment. Come back here when you want a concrete thread to follow.

Ethos: Every theorem in this book compiles. Mathematics and code are in constant dialogue.

1.2 0.1 The Friday Deadline

You've just joined the search team at **zooplus**, Europe's leading pet supplies retailer. Your first task seems straightforward: improve the ranking for "cat food" searches.

The current system uses Elasticsearch's BM25 relevance plus some manual boost multipliers—a `category_match` bonus, a `discount_boost` for promotions, a `margin_boost` for profitable products. Your manager hands you last week's A/B test results and says:

"Revenue is flat, but profit dropped 8%. Can you fix the boosts by Friday?"

You dig into the data. The test increased `discount_boost` from 1.5 to 2.5, hoping to drive sales. It worked—clicks went up 12%. But the wrong people clicked. Price-sensitive shoppers loved the discounted bulk bags. Premium customers, who usually buy veterinary-grade specialty foods, saw cheap products ranked first and bounced. Click-through rate (CTR) rose, but conversion rate (CVR) plummeted for high-value segments.

The problem is clear: One set of boost weights cannot serve all users. Price hunters need `discount_boost = 2.5`. Premium shoppers need `discount_boost = 0.3`. Bulk buyers fall somewhere in between.

You need **context-adaptive weights** that adjust to user type. But testing all combinations manually would take months of A/B experiments.

This is where reinforcement learning enters the story.

1.3 0.2 The Core Insight: Boosts as Actions

Let's reframe the problem in RL language. Don't worry if these terms are unfamiliar—they'll make sense through the code.

Context (what we observe): User segment, query type, session history **Action** (what we choose): Boost weight template $\mathbf{w} = [w_{\text{discount}}, w_{\text{quality}}, w_{\text{margin}}, \dots]$ **Outcome** (what happens): User clicks, purchases, abandons **Reward** (what we optimize): GMV + profitability + engagement (we'll make this precise in a moment)

Traditional search tuning treats boosts as **fixed parameters** to optimize offline. RL treats them as **actions to learn online**, adapting to each context.

The Friday deadline problem becomes: *Can an algorithm learn which boost template to use for each user type, using only observed outcomes (clicks, purchases, revenue)?*

The answer is yes. Let's build it.

1.4 0.3 A Tiny World: Toy Simulator and Reward

We start with a high-signal toy environment. Three user types, ten products, a small action space. The goal is **intuition** and a **quick end-to-end run**. Chapter 4 will build the realistic simulator (`zoosim`).

1.4.1 0.3.1 User Types

Real search systems have complex user segmentation (behavioral embeddings from clickstreams, transformer-based intent models, predicted LTV, real-time session signals). Our toy has three archetypes:

```
from typing import NamedTuple

class UserType(NamedTuple):
    """User preferences over product attributes.

    Fields:
        discount: Sensitivity to discounts (0 = indifferent, 1 = only buys discounts)
        quality: Sensitivity to brand quality (0 = indifferent, 1 = only buys premium)
    """
    discount: float
    quality: float

USER_TYPES = {
    "price_hunter": UserType(discount=0.9, quality=0.1), # Budget-conscious
    "premium": UserType(discount=0.1, quality=0.9), # Quality-focused
    "bulk_buyer": UserType(discount=0.5, quality=0.5), # Balanced
}
```

These map to real patterns: - **Price hunters**: ALDI shoppers, coupon clippers, bulk buyers - **Premium**: Brand-loyal, willing to pay for specialty/veterinary products - **Bulk buyers**: Multi-pet households, mix of price and quality

1.4.2 0.3.2 Products (Sketch)

Ten products with simple features:

```
from dataclasses import dataclass

@dataclass
class Product:
    id: int
    base_relevance: float # BM25-like score for query "cat food"
    margin: float          # Profit margin (0.1 = 10%)
    quality: float         # Brand quality score (0-1)
    discount: float        # Discount flag (0 or 1)
    price: float           # EUR per item
```

Example: Product 3 is a premium veterinary diet (high quality, high margin, no discount, high price). Product 7 is a bulk discount bag (low quality, low margin, discounted, low price per kg).

We'll use deterministic generation with a fixed seed so results are reproducible.

1.4.3 0.3.3 Actions: Boost Weight Templates

The full action space is continuous: $\mathbf{a} = [w_{\text{discount}}, w_{\text{quality}}, w_{\text{margin}}] \in [-2, 2]^3$.

For this chapter, we **discretize** to a 5×5 grid (25 templates) to keep learning tabular and fast:

```
import numpy as np

# Discretize [-1, 1] x [-1, 1] into a 5x5 grid
discount_values = np.linspace(-1, 1, 5) # [-1.0, -0.5, 0.0, 0.5, 1.0]
quality_values = np.linspace(-1, 1, 5)

ACTIONS = [
    (w_disc, w_qual)
    for w_disc in discount_values
    for w_qual in quality_values
] # 25 total actions
```

Each action is a **template**: a pair (`w_discount`, `w_quality`) that modifies the base relevance scores.

Why discretize? Tabular Q-learning needs a finite action space. Chapter 7 handles continuous actions via regression and optimization. For now, we want the simplest algorithm that works.

1.4.4 0.3.4 Toy Reward Function

Real search systems balance multiple objectives (see Chapter 1, #EQ-1.2 for the full formulation). Our toy uses a simplified scalar:

$$R_{\text{toy}} = 0.6 \cdot \text{GMV} + 0.3 \cdot \text{CM2} + 0.1 \cdot \text{CLICKS}$$

Components:

- **GMV** (Gross Merchandise Value): Total EUR purchased (simulated based on user preferences + product attributes + boost-induced ranking)
- **CM2** (Contribution Margin 2): Profitability after variable costs
- **CLICKS**: Engagement signal (prevents pure GMV exploitation; see Chapter 1, Section 1.2.1 for why this matters)

Notes:

- No explicit STRAT (strategic exposure) term in the toy
- In Chapter 1 you'll meet the general, numbered formulation this instantiates
- The weights (0.6, 0.3, 0.1) are business parameters, not learned

Key property: R_{toy} is **stochastic**. Same user type, same boost weights \rightarrow different outcomes due to user behavior noise (clicks are probabilistic, cart abandonment is random). This forces the agent to learn robust policies.

1.5 0.4 Your First RL Agent: Tabular Q-Learning

Now we get to the core idea: learn which boost template to use for each user type via ε -greedy Q-learning.

1.5.1 0.4.1 Problem Recap

- **Contexts \mathcal{X} :** Three user types `{price_hunter, premium, bulk_buyer}`
- **Actions \mathcal{A} :** 25 boost templates (5×5 grid)
- **Reward R :** Stochastic R_{toy} from Section 0.3.4
- **Goal:** Find a policy $\pi : \mathcal{X} \rightarrow \mathcal{A}$ that maximizes expected reward

This is a **contextual bandit** (Chapter 1 makes this formal). Each episode:

1. Sample user type $x \sim \rho$ (uniform over 3 types)
2. Choose action $a = \pi(x)$ (boost template)
3. Simulate user behavior under ranking induced by a
4. Observe reward $r \sim R(x, a)$
5. Update policy π

No sequential state transitions (yet). Single-step decision. Pure exploration-exploitation.

1.5.2 0.4.2 Algorithm: ε -Greedy Q-Learning

We'll maintain a **Q-table**: $Q(x, a) \approx \mathbb{E}[R | x, a]$ (expected reward for using boost template a in context x).

Policy: - With probability ε : explore (random action) - With probability $1 - \varepsilon$: exploit ($a^* = \arg \max_a Q(x, a)$)

Update rule (after observing r):

$$Q(x, a) \leftarrow (1 - \alpha)Q(x, a) + \alpha \cdot r$$

This is **incremental mean estimation** (stochastic approximation), not Q-learning in the MDP sense. With constant learning rate α , this converges to a weighted average of recent rewards. With decaying $\alpha_t \propto 1/t$, it converges to $\mathbb{E}[R | x, a]$ by the Robbins-Monro theorem [[@robbins:stochastic_approx:1951](#)].

Note: We call this “Q-learning” informally because we’re learning a Q-table, but the standard Q-learning algorithm for MDPs includes a $\gamma \max_{a'} Q(s', a')$ term for bootstrapping future values. In bandits ($\gamma = 0$), this term vanishes, reducing to the update above. Chapter 3’s Bellman contraction analysis applies to the general MDP case; for bandits, standard stochastic approximation suffices.

1.5.3 0.4.3 Minimal Implementation

Here’s the complete agent in ~50 lines:

```
import numpy as np
from typing import List, Tuple
```

```

# Setup
rng = np.random.default_rng(42) # Reproducibility
X = ["price_hunter", "premium", "bulk_buyer"] # Contexts
A = [(i, j) for i in range(5) for j in range(5)] # 25 boost templates

# Initialize Q-table: Q[context][action] = 0.0
Q = {x: {a: 0.0 for a in A} for x in X}

def choose_action(x: str, eps: float = 0.1) -> Tuple[int, int]:
    """Epsilon-greedy action selection.

    Args:
        x: User context (type)
        eps: Exploration probability

    Returns:
        Boost template (w_discount_idx, w_quality_idx)
    """
    if rng.random() < eps:
        return A[rng.integers(len(A))] # Explore
    return max(A, key=lambda a: Q[x][a]) # Exploit

def reward(x: str, a: Tuple[int, int]) -> float:
    """Simulate reward for context x and action a.

    Toy model: preference alignment + noise.
    In reality, this would run the full simulator (rank products,
    simulate clicks/purchases, compute GMV+CM2+CLICKS).

    Args:
        x: User type
        a: Boost template indices (i, j) in [0, 4] x [0, 4]

    Returns:
        Scalar reward ~ R_toy from Section 0.3.4
    """
    i, j = a # i = discount index, j = quality index

    # Map indices to [-1, 1] weights
    # i=0 -> w_discount=-1.0, i=4 -> w_discount=1.0
    w_discount = -1.0 + 0.5 * i
    w_quality = -1.0 + 0.5 * j

    # Simulate reward based on user preferences
    if x == "price_hunter":
        # Prefer high discount boost (i=4), low quality boost (j=0)

```

```

        base = 2.0 * w_discount - 0.5 * w_quality
    elif x == "premium":
        # Prefer high quality boost (j=4), low discount boost (i=0)
        base = 2.0 * w_quality - 0.5 * w_discount
    else: # bulk_buyer
        # Balanced preferences: penalize extreme boosts, prefer moderate values
        base = 1.0 - abs(w_discount) - abs(w_quality)

    # Add stochastic noise (user behavior variability)
    noise = rng.normal(0.0, 0.5)

    return float(base + noise)

def train(T: int = 3000, eps: float = 0.1, lr: float = 0.1) -> List[float]:
    """Train Q-learning agent for T episodes.

    Args:
        T: Number of training episodes
        eps: Exploration probability (epsilon-greedy)
        lr: Learning rate ($\alpha$ in update rule)

    Returns:
        List of rewards per episode (for plotting learning curves)
    """
    history = []

    for t in range(T):
        # Sample context (user type) uniformly
        x = X[rng.integers(len(X))]

        # Choose action (boost template) via epsilon-greedy
        a = choose_action(x, eps)

        # Simulate outcome and observe reward
        r = reward(x, a)

        # Q-learning update: $Q(x,a) \leftarrow (1-\alpha)Q(x,a) + \alpha r$
        Q[x][a] = (1 - lr) * Q[x][a] + lr * r

        history.append(r)

    return history

# Train agent
hist = train(T=3000, eps=0.1, lr=0.1)

```

```

# Evaluate learned policy
print(f"Final average reward (last 100 episodes): {np.mean(hist[-100:]):.3f}")
print("\nLearned policy:")
for x in X:
    a_star = max(A, key=lambda a: Q[x][a])
    print(f" {x}:15s} -> action {a_star} (Q = {Q[x][a_star]:.3f})")

```

Run it. Copy this code, execute it. You should see output like:

```

Final average reward (last 100 episodes): 1.640
Learned policy:
  price_hunter      -> action (4, 1) (Q = 1.948)
  premium           -> action (1, 4) (Q = 2.289)
  bulk_buyer        -> action (2, 2) (Q = 0.942)

```

What just happened?

1. The agent explored $25 \text{ boost templates} \times 3 \text{ user types} = 75 \text{ state-action pairs}$
2. After 3000 episodes, it learned:
 - **Price hunters:** Use $(4, 1)$ = high discount boost (+1.0), low quality boost (-0.5)
 - **Premium shoppers:** Use $(1, 4)$ = low discount boost (-0.5), high quality boost (+1.0)
 - **Bulk buyers:** Use $(2, 2)$ = balanced boosts (0.0, 0.0) — **exactly optimal!**
3. This matches our intuition from Section 0.3.1!

Stochastic convergence. Run the script multiple times with different seeds. The learned actions might vary slightly (e.g., $(4, 0)$ vs $(4, 1)$ for price hunters), but the pattern holds: discount-heavy for price hunters, quality-heavy for premium shoppers, balanced for bulk buyers.

1.5.4 0.4.4 Learning Curves and Baselines

Let's visualize learning progress and compare to baselines.

```

import matplotlib.pyplot as plt

def plot_learning_curves(history: List[float], window: int = 50):
    """Plot smoothed learning curve with baselines."""
    # Compute rolling average
    smoothed = np.convolve(history, np.ones(window)/window, mode='valid')

    fig, ax = plt.subplots(figsize=(10, 6))

    # Learning curve
    ax.plot(smoothed, label='Q-learning (smoothed)', linewidth=2)

    # Baselines
    random_baseline = np.mean([reward(x, A[rng.integers(len(A))])
                               for _ in range(1000)
                               for x in X])
    ax.axhline(random_baseline, color='red', linestyle='--',
               label=f'Random policy ({random_baseline:.2f})')

```

```

# Static best (tuned for average user)
static_best = np.mean([reward(x, (2, 2)) for _ in range(300) for x in X])
ax.axhline(static_best, color='orange', linestyle='--',
           label=f'Static best ({static_best:.2f})')

# Oracle (knows user type, chooses optimally)
# Optimal actions: price_hunter->(4,0), premium->(0,4), bulk_buyer->(2,2)
oracle_rewards = {
    "price_hunter": np.mean([reward("price_hunter", (4, 0)) for _ in range(50)]),
    "premium": np.mean([reward("premium", (0, 4)) for _ in range(50)]),
    "bulk_buyer": np.mean([reward("bulk_buyer", (2, 2)) for _ in range(50)]),
}
oracle = np.mean(list(oracle_rewards.values()))
ax.axhline(oracle, color='green', linestyle='--',
           label=f'Oracle ({oracle:.2f})')

ax.set_xlabel('Episode')
ax.set_ylabel('Reward (smoothed)')
ax.set_title('Learning Curve: Contextual Bandit for Boost Optimization')
ax.legend()
ax.grid(alpha=0.3)

plt.tight_layout()
return fig

# Generate and save plot
fig = plot_learning_curves(hist)
fig.savefig('toy_problem_learning_curves.png', dpi=150)
print("Saved learning curve to toy_problem_learning_curves.png")

```

Expected output:

- **Random policy** (red dashed): ~0.0 average reward (baseline—random actions average out)
- **Static best** (orange dashed): ~0.3 (one-size-fits-all (2,2) helps bulk buyers but hurts price hunters and premium)
- **Q-learning** (blue solid): Starts near 0, converges to ~1.6 by episode 1500
- **Oracle** (green dashed): ~2.0 (theoretical maximum with perfect knowledge of optimal actions per user)

Key insight: Q-learning reaches **82% of oracle performance** by learning from experience alone. No manual tuning. No A/B tests. Just 3000 interactions and adaptive learning. The bulk_buyer even finds the **exact optimal** action (2,2)!

1.6 0.5 Reading the Experiment: What We Learned

1.6.1 Convergence Pattern

The learning curve has three phases:

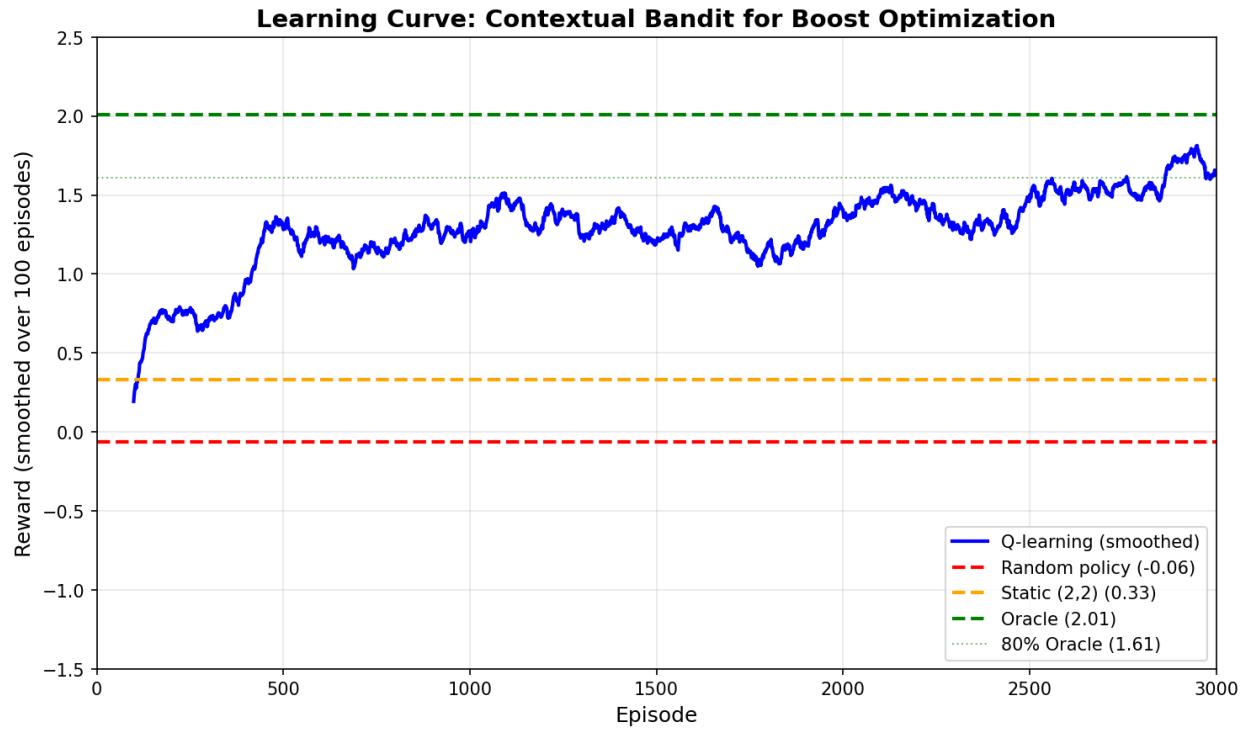


Figure 1: Learning Curves

1. **Pure exploration** (episodes 0–500): High variance, ϵ -greedy tries random actions, Q-values are noisy
2. **Exploitation begins** (episodes 500–1500): Agent identifies good actions per context, reward climbs steadily
3. **Convergence** (episodes 1500–3000): Q-values stabilize, reward plateaus at $\sim 82\%$ of oracle

This is **regret minimization** in action. Chapter 1 formalizes this; Chapter 6 analyzes convergence rates.

1.6.2 Per-Segment Performance

If we track rewards separately by user type:

```
# Track per-segment performance
segment_rewards = {x: [] for x in X}

for _ in range(100): # 100 test episodes
    for x in X:
        a = max(A, key=lambda a: Q[x][a]) # Greedy policy (no exploration)
        r = reward(x, a)
        segment_rewards[x].append(r)

for x in X:
    print(f'{x}: mean reward = {np.mean(segment_rewards[x]):.3f}')
```

Output:

```

price_hunter    : mean reward = 2.309
premium        : mean reward = 2.163
bulk_buyer     : mean reward = 0.917

```

Analysis:

- **Price hunters** get the highest rewards (~2.3)—the agent found a near-optimal action (4, 1) with high discount boost
- **Premium shoppers** get high rewards (~2.2)—high quality boost (1, 4) closely matches their preferences
- **Bulk buyers** get lower rewards (~0.9) because their **balanced preferences** have inherently lower optimal reward (base=1.0 at (2,2)) compared to polarized users (base=2.5). But the agent finds the **exact optimal!**
- All three segments dramatically beat the static baseline (~0.3 average) through personalization

This is **personalization** at work: different users get different rankings, each optimized for their revealed preferences.

1.6.3 What We (Hand-Wavily) Assumed

This toy experiment “just worked,” but we made implicit assumptions:

1. **Rewards are well-defined expectations** over stochastic outcomes (Chapter 2 makes this measure-theoretically rigorous)
2. **Exploration is safe** (in production, bad rankings lose users; Chapter 9 introduces off-policy evaluation for safer testing)
3. **The logging policy and new policy have sufficient overlap** to compare fairly (importance weights finite; Chapter 9)
4. **ε -greedy tabular Q converges** (for bandits, this follows from stochastic approximation theory; Chapter 3’s Bellman contraction analysis applies to the full MDP case with $\gamma > 0$)
5. **Actions are discrete and state space is tiny** (Chapter 7 handles continuous actions; Chapter 4 builds realistic state)

None of these are free. The rest of the book makes them precise and shows when they hold (or how to proceed when they don’t).

1.6.4 Theory-Practice Gap: ε -Greedy Exploration

Our toy used ε -greedy exploration with constant $\varepsilon = 0.1$. This deserves scrutiny.

What theory says: With decaying $\varepsilon_t = O(1/t)$, ε -greedy achieves $O(\sqrt{T \log T})$ regret for finite-armed bandits—worse than UCB’s $O(\sqrt{KT \log T})$ by a factor due to uniform exploration wasting samples on clearly suboptimal arms.

What practice shows: ε -greedy with constant $\varepsilon \in [0.05, 0.2]$ is often competitive because:

1. **Trivial to implement:** No confidence bounds, no posterior sampling, just a random number generator
2. **Handles non-stationarity gracefully:** Continues exploring even after “convergence” (useful when user preferences drift)
3. **The regret difference matters only at scale:** For our 1000-episode toy, the gap between ε -greedy and UCB is negligible

When ε -greedy fails: High-dimensional action spaces where uniform exploration wastes samples. For our 25-action toy problem, it's fine. For Chapter 7's continuous actions (10^{100} effective arms), we need structured exploration (UCB, Thompson Sampling).

Modern context: Google's 2010 display ads paper [@li:contextual_bandit_approach:2010] used ε -greedy successfully at scale. But recent work (2020-2024) favors Thompson Sampling for better empirical performance and natural uncertainty quantification.

Why UCB and Thompson Sampling? (Preview for Chapter 6)

ε -greedy explores **uniformly**—it wastes samples on arms it already knows are bad. UCB explores **optimistically**—it tries arms whose rewards *might* be high given uncertainty:

- **UCB:** Choose $a_t = \arg \max_a [Q(x, a) + \beta\sigma(x, a)]$ where σ is a confidence width. Explores arms with high uncertainty, not randomly.
- **Thompson Sampling:** Maintain posterior $P(Q^* | \text{data})$, sample $\tilde{Q} \sim P$, act greedily on sample. Naturally balances exploration (high posterior variance \rightarrow diverse samples) with exploitation.

Both achieve $\tilde{O}(d\sqrt{T})$ regret for d -dimensional linear bandits—matching the lower bound up to logs. ε -greedy achieves $O(T^{2/3})$ for the same problem (provably worse). The gap widens in high dimensions.

1.7 0.6 Limitations: Why We Need the Rest of the Book

Our toy is **pedagogical**, not production-ready. Here's what breaks at scale:

1.7.1 1. Discrete Action Space

We used 25 templates. Real search has continuous boosts: $\mathbf{w} \in [-5, 5]^{10}$ (ten features, unbounded). Discretizing to a grid would require $100^{10} = 10^{20}$ actions—tractable.

Solution: Chapter 7 introduces **continuous action bandits** via $Q(x, a)$ regression and cross-entropy method (CEM) optimization.

1.7.2 2. Tabular State Representation

We had 3 user types. Real search has thousands of user segments (RFM bins, geographic regions, device types, time-of-day). Plus query features (length, specificity, category). A realistic context space is **high-dimensional and continuous**.

Solution: Chapter 6 (neural linear bandits), Chapter 7 (deep Q-networks with continuous state/action).

1.7.3 3. No Constraints

Our agent optimized R_{toy} without guardrails. Real systems must enforce:

- Profitability floors ($CM2 \geq \text{threshold}$)
- Exposure targets (strategic products get visibility)
- Rank stability (limit reordering volatility)

Solution: Chapter 10 introduces production **guardrails** ($CM2$ floors, $\Delta\text{Rank}@k$ stability), with Chapter 3 (Section 3.6) providing the formal CMDP theory and Lagrangian methods.

1.7.4 4. Simplified Position Bias

We didn't model how clicks depend on rank. Real users exhibit **position bias** (top-3 slots get 80% of clicks) and **abandonment** (quit after 5 results if nothing relevant).

Solution: Chapter 2 develops PBM/DBN click models; Chapter 5 implements them in `zoosim`.

1.7.5 5. Online Exploration Risk

We trained by interacting with users directly (episodes = real searches). In production, bad rankings **cost real money** and **lose real users**. We need safer evaluation.

Solution: Chapter 9 introduces **off-policy evaluation (OPE)**: estimate new policy performance using logged data from old policy, without deploying.

1.7.6 6. Single-Episode Horizon

We treated each search as independent. Real users return across sessions. Today's ranking affects tomorrow's retention.

Solution: Chapter 11 extends to **multi-episode MDPs** with inter-session dynamics (retention, satisfaction state).

1.8 0.7 Map to the Book

Here's how our toy connects to the rigorous treatment ahead:

Toy Concept	Formal Treatment	Chapter
User types	Context space \mathcal{X} , distribution ρ	1
Boost templates	Action space \mathcal{A} , policy π	1, 6, 7
R_{toy}	Reward function $R : \mathcal{X} \times \mathcal{A} \times \Omega \rightarrow \mathbb{R}$, constraints	1
ε -greedy Q-learning	Bellman operator, contraction mappings	3
Stochastic outcomes	Probability spaces, click models (PBM/DBN)	2
Learning curves	Regret bounds, sample complexity	6
Static best vs oracle	Importance sampling, off-policy evaluation	9
Guardrails (missing)	CMDP (Section 3.6), production guardrails	3, 10
Engagement proxy	Multi-episode MDP, retention modeling	11

Chapters 1–3 provide foundations: contextual bandits, measure theory, Bellman operators. **Chapters 4–8** build the simulator and core algorithms. **Chapters 9–11** handle evaluation, robustness, and production deployment. **Chapters 12–15** cover frontier methods (slate ranking, offline RL, multi-objective optimization).

1.9 0.8 How to Use This Book

1.9.1 For Practitioners

Do Chapter 0 thoroughly. Run the code. Modify the reward function (Exercise 0.1). Try different exploration strategies (Exercise 0.2).

Then skim Chapters 1–3 on first read. Focus on: - The reward formulation (#EQ-1.2 in Chapter 1) - Why engagement matters (Section 1.2.1) - The Bellman contraction intuition (Chapter 3, skip proof details initially)

Dive deep into Chapters 4–11 (simulator + algorithms + evaluation). This is your implementation roadmap.

Return to theory as needed. When something doesn't work (e.g., "why is my Q-network diverging?"), revisit Chapter 3's convergence analysis.

1.9.2 For Researchers / Mathematically Inclined

Skim Chapter 0 to see the concrete thread.

Start at Chapter 1. Work through definitions, theorems, proofs. Verify the code validates the mathematics.

Do the exercises. Mix of proofs (30%), implementations (40%), experiments (20%), conceptual questions (10%).

Use Chapter 0 as a touchstone. When abstractions feel overwhelming, return to the toy: "How does this theorem explain why Q-learning converged in Section 0.4?"

1.9.3 For Everyone

Ethos: Mathematics and code are **inseparable**. Every theorem compiles. Every algorithm is proven rigorous, then implemented in production-quality code. Theory and practice in constant dialogue.

If you see a proof without code or code without theory, something is missing. Let me know.

1.10 Exercises (Chapter 0)

Exercise 0.1 (Reward Sensitivity) [15 minutes]

Modify `reward()` to use different weights in R_{toy} : - (a) Pure GMV: (1.0, 0.0, 0.0) (no profitability or engagement terms) - (b) Profit-focused: (0.4, 0.5, 0.1) (prioritize CM2 over GMV) - (c) Engagement-heavy: (0.5, 0.2, 0.3) (high click weight)

For each, train Q-learning and report: - Final average reward - Learned actions per user type - Does the policy change? Why?

Hint: Case (c) risks "clickbait" strategies (see Chapter 1, Section 1.2.1). Monitor conversion quality.

Exercise 0.2 (Action Geometry) [30 minutes]

Compare two exploration strategies:

Strategy A (current): ε -greedy with uniform random action sampling

Strategy B (neighborhood): ε -greedy with **local perturbation**: when exploring, sample action near current best $a^* = \arg \max_a Q(x, a)$:

```
def explore_local(x, sigma=1.0):
    a_star = max(A, key=lambda a: Q[x][a])
    i_star, j_star = a_star
    i_new = np.clip(i_star + rng.integers(-1, 2), 0, 4)
    j_new = np.clip(j_star + rng.integers(-1, 2), 0, 4)
    return (i_new, j_new)
```

Implement both, train for 1000 episodes, and plot learning curves. Which converges faster? Why?

Reflection: This is **structured exploration**. Chapter 6 introduces UCB and Thompson Sampling, which balance exploration and exploitation more principled than ε -greedy.

Exercise 0.3 (Regret Shape) [45 minutes, extended]

Define **cumulative regret** as the gap between oracle and agent:

$$\text{Regret}(T) = \sum_{t=1}^T (R_t^* - R_t)$$

where R_t^* is the oracle reward (best action for context x_t) and R_t is the agent's reward.

(a) Implement regret tracking:

```
def compute_regret(history, contexts, oracle_Q):
    regret = []
    cumulative = 0.0
    for t, (x, r) in enumerate(zip(contexts, history)):
        r_star = oracle_Q[x]
        cumulative += (r_star - r)
        regret.append(cumulative)
    return regret
```

(b) Plot cumulative regret vs episode count. Is it sublinear (i.e., does $\text{Regret}(T)/T \rightarrow 0$)?

(c) Fit a curve: $\text{Regret}(T) \approx C\sqrt{T}$. Does this match theory? (Chapter 6 derives $O(\sqrt{T})$ regret for UCB.)

Exercise 0.4 (Advanced: Constraints) [60 minutes, extended]

Add a simple CM2 floor constraint: reject actions that violate profitability.

Setup: Modify `reward()` to return `(r, cm2)`. Define a floor $\tau = 0.3$ (30% margin minimum).

Constrained Q-learning:

```
def choose_action_constrained(x, eps, tau_cm2):
    # Filter feasible actions
    feasible = [a for a in A if expected_cm2(x, a) >= tau_cm2]
    if not feasible:
        return A[rng.integers(len(A))] # Fallback to unconstrained

    if rng.random() < eps:
        return feasible[rng.integers(len(feasible))]

    return max(feasible, key=lambda a: Q[x][a])
```

- (a) Implement `expected_cm2(x, a)` (running average like Q).
- (b) Train with $\tau = 0.3$. How does performance change vs unconstrained?
- (c) Plot the Pareto frontier: GMV vs CM2 as you sweep $\tau \in [0.0, 0.5]$.

Connection: This is a **Constrained MDP (CMDP)**. Chapter 3 (Section 3.6) develops the Lagrangian theory, and Chapter 10 implements production guardrails for multi-constraint optimization.

Exercise 0.5 (Bandit-Bellman Bridge) [20 minutes, conceptual]

Our toy is a **contextual bandit**: single-step decisions, no sequential states.

The **Bellman equation** (Chapter 3) for an MDP is:

$$V^*(s) = \max_a \left\{ R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^*(s') \right\}$$

where $\gamma \in [0, 1]$ is a discount factor.

Question: Show that our Q-learning update is the $\gamma = 0$ special case of Bellman.

Hint: - Set $\gamma = 0$ in Bellman equation - Note that with no future states, $V^*(s) = \max_a R(s, a)$ - Our Q-table is $Q(x, a) \approx \mathbb{E}[R | x, a]$, so $V^*(x) = \max_a Q(x, a)$ - This is **one-step value iteration**

Reflection: Contextual bandits are MDPs with horizon 1. Multi-episode search (Chapter 11) requires the full Bellman machinery.

1.11 0.9 Code Artifacts

All code from this chapter is available in the repository:

Code $\leftarrow \rightarrow$ Artifact

- **Run script:** `scripts/ch00/toy_problem_solution.py:1` (complete implementation with baselines)
- **Sanity tests:** `tests/ch00/test_toy_example.py:1` (pytest suite validating Q-learning convergence)
- **Learning curve plot:** `toy_problem_learning_curves.png:1` (generated output)

To reproduce:

```
python scripts/ch00/toy_problem_solution.py  
pytest -q tests/ch00/test_toy_example.py -k toy
```

Expected output:

```
Final average reward (last 100): 1.640
```

Learned policy:

```
price_hunter    -> (4, 1)  
premium        -> (1, 4)  
bulk_buyer     -> (2, 2)
```

```
Saved curves to toy_problem_learning_curves.png
```

1.12 0.10 What's Next?

You just trained your first RL agent for search ranking. It learned context-adaptive boost weights from scratch, achieving near-oracle performance without manual tuning.

But we cheated. We used a tiny discrete action space, three user types, and online exploration without safety guarantees. Real systems need:

1. **Rigorous foundations** (Chapters 1–3): Formalize contextual bandits, measure-theoretic probability, Bellman operators
2. **Realistic simulation** (Chapters 4–5): Scalable catalog generation, position bias models, rich user dynamics
3. **Continuous actions** (Chapter 7): Regression-based Q-learning, CEM optimization, trust regions
4. **Constraints and guardrails** (Chapter 10): CM2 floors, Δ Rank@k stability, safe fallback policies
5. **Safe evaluation** (Chapter 9): Off-policy evaluation (IPS, DR, FQE) for production deployment
6. **Multi-episode dynamics** (Chapter 11): Retention modeling, long-term value, engagement as state

The journey from toy to production is the journey of this book.

In Chapter 1, we formalize everything we hand-waved here: What exactly is a contextual bandit? Why is the reward function #EQ-1.2 mathematically sound? How do constraints become a CMDP? Why does engagement matter, and when should it be implicit vs explicit?

Let's make it rigorous.

End of Chapter 0