

## Contents

<b>1 Chapter 6 Winning Parameters: LinUCB vs Thompson Sampling</b>	<b>1</b>
1.1 Configuration A: Rich Features + Oracle Latents (LinUCB Wins) . . . . .	1
1.1.1 Expected Results (Oracle Latents) . . . . .	1
1.2 Configuration B: Rich Features + Estimated Latents (TS Wins) . . . . .	2
1.2.1 Expected Results (Estimated Latents) . . . . .	2
1.3 The Algorithm Selection Principle . . . . .	2
1.4 Hyperparameter Reference . . . . .	3
1.5 Running the Full Compute Arc . . . . .	3
1.6 References . . . . .	3

## 1 Chapter 6 Winning Parameters: LinUCB vs Thompson Sampling

This document records the **Algorithm Selection Principle** discovered in Chapter 6:

**Algorithm selection depends on feature quality.** - Clean/Oracle features → LinUCB (precise exploitation) - Noisy/Estimated features → Thompson Sampling (robust exploration)

### 1.1 Configuration A: Rich Features + Oracle Latents (LinUCB Wins)

When you have access to true user latent preferences (oracle features), LinUCB's precise exploitation wins:

```
python scripts/ch06/template_bandits_demo.py \
--features rich \
--rich-regularization blend \
--n-static 2000 \
--n-bandit 20000 \
--world-seed 20250322 \
--bandit-base-seed 20250349 \
--prior-weight 50 \
--lin-alpha 0.2 \
--ts-sigma 0.5
```

#### 1.1.1 Expected Results (Oracle Latents)

Policy	GMV	ΔGMV vs Static
Static-Premium	7.11	baseline
<b>LinUCB</b>	<b>9.42</b>	<b>+32.5%</b>
Thompson Sampling	9.39	+32.1%

**Winner:** LinUCB (marginally, +0.4 percentage points)

**Why LinUCB wins:** With clean oracle features, the linear model assumption holds nearly perfectly. LinUCB's UCB exploration bonus shrinks precisely as uncertainty decreases, enabling pre-

cise exploitation. Note that TS also performs excellently with oracle features—the real differentiation emerges under feature noise (see Estimated Latents below).

---

## 1.2 Configuration B: Rich Features + Estimated Latents (TS Wins)

In production, you don't have oracle latents—you must estimate user preferences from behavior. With noisy estimated features, Thompson Sampling's robust exploration wins:

```
python scripts/ch06/template_bandits_demo.py \
    --features rich_est \
    --n-static 2000 \
    --n-bandit 20000 \
    --world-seed 20250322 \
    --bandit-base-seed 20250349 \
    --prior-weight 50 \
    --lin-alpha 0.2 \
    --ts-sigma 0.5
```

### 1.2.1 Expected Results (Estimated Latents)

Policy	GMV	$\Delta$ GMV vs Static
Static-Premium	7.11	baseline
LinUCB	7.52	+5.8%
<b>Thompson Sampling</b>	<b>9.31</b>	<b>+31.0%</b>

**Winner:** Thompson Sampling by 25 percentage points

**Why TS wins:** With noisy estimated features, LinUCB can lock into suboptimal templates based on spurious correlations. Thompson Sampling's posterior sampling maintains perpetual exploration variance, hedging against feature noise and avoiding premature convergence.

---

## 1.3 The Algorithm Selection Principle

Feature Quality	Winner	Margin	Recommendation
Oracle (clean)	LinUCB	+0.4 pts	Use when features are direct measurements or carefully validated
Estimated (noisy)	TS	+25 pts	<b>Default for production</b>

**Key insight:** With oracle features, both algorithms perform nearly identically (LinUCB 9.42, TS 9.39). The dramatic differentiation emerges **only under feature noise**—TS maintains +31% lift while LinUCB drops to +6%.

**Production implication:** Real e-commerce systems have noisy estimated features— inferred from clicks, aggregated from proxies, delayed from behavior logs. **Default to Thompson Sampling in production.**

---

## 1.4 Hyperparameter Reference

Parameter	Value	Notes
--prior-weight	50	Pseudo-count for static priors
--lin-alpha	0.2	LinUCB exploration bonus (conservative)
--ts-sigma	0.5	Thompson Sampling posterior std
--rich-regularization	blend	Regularization mode for rich features
--n-static	2000	Episodes for static baseline evaluation
--n-bandit	20000	Episodes for bandit training

---

## 1.5 Running the Full Compute Arc

To reproduce the complete three-stage narrative:

```
python scripts/ch06/ch06_compute_arc.py \
    --n-static 2000 \
    --n-bandit 20000 \
    --base-seed 20250322 \
    --out-dir docs/book/drafts/ch06/data \
    --prior-weight 50 \
    --lin-alpha 0.2 \
    --ts-sigma 0.5
```

This runs all three stages: 1. **Simple features** (failure mode) 2. **Rich features + oracle latents** (LinUCB wins) 3. **Rich features + estimated latents** (TS wins)

And produces JSON summaries + comparison tables demonstrating the Algorithm Selection Principle.

---

## 1.6 References

- Chapter 6 main narrative: [docs/book/drafts/ch06/discrete\\_template\\_bandits.md](#)
- Lab exercises: [docs/book/drafts/ch06/exercises\\_labs.md](#)
- Compute arc script: [scripts/ch06/ch06\\_compute\\_arc.py](#)