

Chapter 0: Motivation

Contents

1 Chapter 0 — Motivation: A First RL Experiment	1
1.1 0.0 Who Should Read This?	1
1.2 0.1 The Friday Deadline	1
1.3 0.2 The Core Insight: Boosts as Actions	2
1.4 0.3 A Tiny World: Toy Simulator and Reward	2
1.5 0.4 A First RL Agent: Tabular Q-Learning	4
1.6 0.5 Reading the Experiment: What We Learned	9
1.7 0.6 Limitations: Why We Need the Rest of the Book	10
1.8 0.7 Map to the Book	11
1.9 0.8 How to Use This Book	12
1.10 Exercises (Chapter 0)	12
1.11 0.9 Code Artifacts	14
1.12 0.10 What's Next?	15

1 Chapter 0 — Motivation: A First RL Experiment

Vlad Prytula

1.1 0.0 Who Should Read This?

This chapter is an optional warm-up: it is deliberately light on mathematics and heavy on code. We build a tiny search world, train a small agent to learn context-adaptive boost weights, and observe the core RL loop in action. Chapters 1–3 provide the rigorous foundations that explain *why* the experiment works and *when* it fails.

Two reading paths work well:

- Practitioner track: begin here; the goal is a working end-to-end system in ~30 minutes, then return to theory as needed.
- Foundations track: skim this chapter for the concrete thread, then begin Chapter 1’s rigorous development; we return here whenever an example is useful.

Ethos: every theorem in this book compiles. Mathematics and code are in constant dialogue.

1.2 0.1 The Friday Deadline

We consider the following scenario. We have just joined the search team at zooplus, Europe’s leading pet supplies retailer. Our first task seems straightforward: improve the ranking for “cat food” searches.

The current system uses Elasticsearch’s BM25 relevance plus some manual boost multipliers—a `category_match` bonus, a `discount_boost` for promotions, a `margin_boost` for profitable products. Our manager hands us last week’s A/B test results and says:

“Revenue is flat, but profit dropped 8%. Can we fix the boosts by Friday?”

We dig into the data. The test increased `discount_boost` from 1.5 to 2.5, hoping to drive sales. It worked—clicks went up 12%. But the wrong people clicked. Price-sensitive shoppers loved the discounted bulk bags. Premium customers, who usually buy veterinary-grade specialty foods, saw cheap products ranked first and bounced. Click-through rate (CTR) rose, but conversion rate (CVR) plummeted for high-value segments.

The problem is clear: one set of boost weights cannot serve all users. Price hunters need `discount_boost = 2.5`. Premium shoppers need `discount_boost = 0.3`. Bulk buyers fall somewhere in between.

We need **context-adaptive weights** that adjust to user type. But testing all combinations manually would take months of A/B experiments.

This is where reinforcement learning enters the story.

1.3 0.2 The Core Insight: Boosts as Actions

We reframe the problem in RL language. If any terminology is unfamiliar, we treat it as a working placeholder: Chapters 1–3 make each object precise and state the assumptions under which it is well-defined.

Context (what we observe): User segment, query type, session history **Action** (what we choose): Boost weight template $\mathbf{w} = [w_{\text{discount}}, w_{\text{quality}}, w_{\text{margin}}, \dots]$ **Outcome** (what happens): User clicks, purchases, abandons **Reward** (what we optimize): GMV + profitability + engagement (we’ll make this precise in a moment)

Traditional search tuning treats boosts as **fixed parameters** to optimize offline. RL treats them as **actions to learn online**, adapting to each context.

The Friday deadline problem becomes: *Can an algorithm learn which boost template to use for each user type, using only observed outcomes (clicks, purchases, revenue)?*

The answer is yes; we now build it.

1.4 0.3 A Tiny World: Toy Simulator and Reward

We start with a high-signal toy environment. Three user types, ten products, a small action space. The goal is intuition and a quick end-to-end run. Chapter 4 builds the realistic simulator (`zoosim`).

1.4.1 0.3.1 User Types

Real search systems have complex user segmentation (behavioral embeddings from clickstreams, transformer-based intent models, predicted LTV, real-time session signals). Our toy has three archetypes:

```
from typing import NamedTuple

class UserType(NamedTuple):
    """User preferences over product attributes.

Fields:
    discount: Sensitivity to discounts (0 = indifferent, 1 = only buys discounts)
    quality: Sensitivity to brand quality (0 = indifferent, 1 = only buys premium)
    """
    discount: float
    quality: float

USER_TYPES = {
    "price_hunter": UserType(discount=0.9, quality=0.1),  # Budget-conscious
```

```

    "premium":      UserType(discount=0.1, quality=0.9),  # Quality-focused
    "bulk_buyer":   UserType(discount=0.5, quality=0.5),  # Balanced
}

```

These map to real patterns: - **Price hunters**: ALDI shoppers, coupon clippers, bulk buyers - **Premium**: Brand-loyal, willing to pay for specialty/veterinary products - **Bulk buyers**: Multi-pet households, mix of price and quality

1.4.2 0.3.2 Products (Sketch)

Ten products with simple features:

```

from dataclasses import dataclass

@dataclass
class Product:
    id: int
    base_relevance: float  # BM25-like score for query "cat food"
    margin: float          # Profit margin (0.1 = 10%)
    quality: float         # Brand quality score (0-1)
    discount: float        # Discount flag (0 or 1)
    price: float           # EUR per item

```

Example: Product 3 is a premium veterinary diet (high quality, high margin, no discount, high price). Product 7 is a bulk discount bag (low quality, low margin, discounted, low price per kg).

We'll use deterministic generation with a fixed seed so results are reproducible.

1.4.3 0.3.3 Actions: Boost Weight Templates

The full action space is continuous: $\mathbf{a} = [w_{\text{discount}}, w_{\text{quality}}, w_{\text{margin}}] \in [-2, 2]^3$.

For this chapter, we **discretize** to a 5×5 grid (25 templates) to keep learning tabular and fast:

```

import numpy as np

# Discretize [-1, 1] x [-1, 1] into a 5x5 grid
discount_values = np.linspace(-1, 1, 5)  # [-1.0, -0.5, 0.0, 0.5, 1.0]
quality_values = np.linspace(-1, 1, 5)

ACTIONS = [
    (w_disc, w_qual)
    for w_disc in discount_values
    for w_qual in quality_values
] # 25 total actions

```

Each action is a **template**: a pair $(w_{\text{discount}}, w_{\text{quality}})$ that modifies the base relevance scores.

Why do we discretize? Tabular Q-learning needs a finite action space. Chapter 7 handles continuous actions via regression and optimization. Here we use the simplest algorithm that works end-to-end.

1.4.4 0.3.4 Toy Reward Function

Real search systems balance multiple objectives (see Chapter 1, EQ-1.2 for the full formulation). Our toy uses a simplified scalar:

$$R_{\text{toy}} = 0.6 \cdot \text{GMV} + 0.3 \cdot \text{CM2} + 0.1 \cdot \text{CLICKS}$$

Components:

- **GMV** (Gross Merchandise Value): Total EUR purchased (simulated based on user preferences + product attributes + boost-induced ranking)
- **CM2** (Contribution Margin 2): Profitability after variable costs
- **CLICKS**: Engagement signal (prevents pure GMV exploitation; see Chapter 1, Section 1.2.1 for why this matters)

Notes:

- No explicit STRAT (strategic exposure) term in the toy
- Chapter 1 presents the general, numbered formulation that this toy instantiates
- The weights (0.6, 0.3, 0.1) are business parameters, not learned

Pedagogical Simplification

The full R_{toy} formula requires simulating user interactions (clicks, purchases, cart dynamics). For Chapter 0's Q-learning demonstration, we use a **closed-form surrogate** (Section 0.4.3) that captures the essential preference-alignment structure without simulator complexity. The true GMV/CM2/click-based reward appears in Chapter 4+ with the full `zoosim` environment.

Key property: R_{toy} is stochastic. The same user type and boost weights can yield different outcomes due to user behavior noise (clicks are probabilistic, cart abandonment is random). This forces the agent to learn robust policies.

1.5 0.4 A First RL Agent: Tabular Q-Learning

We now arrive at the core idea: learn which boost template to use for each user type via ε -greedy tabular learning.

1.5.1 0.4.1 Problem Recap

- **Contexts \mathcal{X}** : Three user types `{price_hunter, premium, bulk_buyer}`
- **Actions \mathcal{A}** : 25 boost templates (5×5 grid)
- **Reward R** : Stochastic R_{toy} from Section 0.3.4
- **Goal**: Find a policy $\pi : \mathcal{X} \rightarrow \mathcal{A}$ that maximizes expected reward

This is a **contextual bandit** (Chapter 1 makes this formal). Each episode:

1. Sample user type $x \sim \rho$ (uniform over 3 types)
2. Choose action $a = \pi(x)$ (boost template)
3. Simulate user behavior under ranking induced by a
4. Observe reward $r \sim R(x, a)$
5. Update policy π

No sequential state transitions (yet). Single-step decision. Pure exploration-exploitation.

1.5.2 0.4.2 Algorithm: ε -Greedy Q-Learning

We'll maintain a **Q-table**: $Q(x, a) \approx \mathbb{E}[R | x, a]$ (expected reward for using boost template a in context x).

Policy: - With probability ε : explore (random action) - With probability $1-\varepsilon$: exploit ($a^* = \arg \max_a Q(x, a)$)

Update rule (after observing r):

$$Q(x, a) \leftarrow (1 - \alpha)Q(x, a) + \alpha \cdot r$$

This is **incremental mean estimation** (stochastic approximation), not Q-learning in the MDP sense. With constant learning rate α , this converges to a weighted average of recent rewards. With decaying $\alpha_t \propto 1/t$, it converges to $\mathbb{E}[R | x, a]$ by the Robbins-Monro theorem (Robbins and Monro 1951).

We call this “Q-learning” informally because we are learning a Q-table, but the standard Q-learning algorithm for MDPs includes a $\gamma \max_{a'} Q(s', a')$ term for bootstrapping future values. In bandits ($\gamma = 0$), this term vanishes, reducing to the update above. Chapter 3’s Bellman contraction analysis applies to the general MDP case; for bandits, standard stochastic approximation suffices.

1.5.3 0.4.3 Minimal Implementation

Here’s the complete agent in ~50 lines.

Pedagogical reward model. Rather than simulate full user interactions (GMV, CM2, clicks), we use a closed-form reward that encodes user preferences directly:

- Price hunters prefer high discount weight (w_{disc})
- Premium users prefer high quality weight (w_{qual})
- Bulk buyers prefer balanced, moderate weights

This surrogate enables rapid Q-learning iterations while preserving the essential optimization structure. The output rewards are **preference-alignment scores** (not EUR), with values typically in $[-2, 3]$.

Discretization note. We index the 5×5 grid of weight pairs for compactness: action (i, j) maps to weights via $w = -1 + 0.5 \cdot \text{index}$. Thus $(0, 0) \mapsto (-1, -1)$ and $(4, 4) \mapsto (1, 1)$.

```
import numpy as np
from typing import List, Tuple

# Setup
rng = np.random.default_rng(42) # Reproducibility
X = ["price_hunter", "premium", "bulk_buyer"] # Contexts
A = [(i, j) for i in range(5) for j in range(5)] # 25 boost templates (indexed)

# Initialize Q-table: Q[context][action] = 0.0
Q = {x: {a: 0.0 for a in A} for x in X}

def choose_action(x: str, eps: float = 0.1) -> Tuple[int, int]:
    """Epsilon-greedy action selection.

    Args:
        x: User context (type)
        eps: Exploration probability

    Returns:
        Boost template (w_discount_idx, w_quality_idx)
    """
    if rng.random() < eps:
        return A[rng.integers(len(A))] # Explore
    return max(A, key=lambda a: Q[x][a]) # Exploit

def reward(x: str, a: Tuple[int, int]) -> float:
    """Simulate reward for context x and action a.

    Toy model: preference alignment + noise.
    In reality, this would run the full simulator (rank products,
    simulate clicks/purchases, compute GMV+CM2+CLICKS).

    Args:
    """

```

```

x: User type
a: Boost template indices (i, j) in [0, 4] x [0, 4]

>Returns:
    Scalar reward ~ R_toy from Section 0.3.4
"""
i, j = a # i = discount index, j = quality index

# Map indices to [-1, 1] weights
# i=0 -> w_discount=-1.0, i=4 -> w_discount=1.0
w_discount = -1.0 + 0.5 * i
w_quality = -1.0 + 0.5 * j

# Simulate reward based on user preferences
if x == "price_hunter":
    # Prefer high discount boost (i=4), low quality boost (j=0)
    base = 2.0 * w_discount - 0.5 * w_quality
elif x == "premium":
    # Prefer high quality boost (j=4), low discount boost (i=0)
    base = 2.0 * w_quality - 0.5 * w_discount
else: # bulk_buyer
    # Balanced preferences: penalize extreme boosts, prefer moderate values
    base = 1.0 - abs(w_discount) - abs(w_quality)

# Add stochastic noise (user behavior variability)
noise = rng.normal(0.0, 0.5)

return float(base + noise)

def train(T: int = 3000, eps: float = 0.1, lr: float = 0.1) -> List[float]:
    """Train Q-learning agent for T episodes.

Args:
    T: Number of training episodes
    eps: Exploration probability (epsilon-greedy)
    lr: Learning rate ($\alpha$ in update rule)

>Returns:
    List of rewards per episode (for plotting learning curves)
"""
history = []

for t in range(T):
    # Sample context (user type) uniformly
    x = X[rng.integers(len(X))]

    # Choose action (boost template) via epsilon-greedy
    a = choose_action(x, eps)

    # Simulate outcome and observe reward
    r = reward(x, a)

    # Q-learning update: Q(x,a) <- (1-alpha)Q(x,a) + alpha*r

```

```

Q[x][a] = (1 - lr) * Q[x][a] + lr * r

history.append(r)

return history

# Train agent
hist = train(T=3000, eps=0.1, lr=0.1)

# Evaluate learned policy
print(f"Final average reward (last 100 episodes): {np.mean(hist[-100:]):.3f}")
print("\nLearned policy:")
for x in X:
    a_star = max(A, key=lambda a: Q[x][a])
    print(f" {x}:15s} -> action {a_star} (Q = {Q[x][a_star]:.3f})")

```

With a fixed seed, we obtain representative output of the form (preference-alignment scores, not EUR):

```

Final average reward (last 100 episodes): 1.640 # preference-alignment scale
Learned policy:
  price_hunter      -> action (4, 1) (Q = 1.948)
  premium           -> action (1, 4) (Q = 2.289)
  bulk_buyer        -> action (2, 2) (Q = 0.942)

```

What just happened?

1. The agent explored $25 \text{ boost templates} \times 3 \text{ user types} = 75$ state-action pairs
2. After 3000 episodes, it learned:
 - **Price hunters:** Use $(4, 1)$ = high discount boost (+1.0), low quality boost (-0.5)
 - **Premium shoppers:** Use $(1, 4)$ = low discount boost (-0.5), high quality boost (+1.0)
 - **Bulk buyers:** Use $(2, 2)$ = balanced boosts (0.0, 0.0) — exactly optimal.
3. This matches our intuition from Section 0.3.1!

Stochastic convergence. Across random seeds, the learned actions might vary slightly (e.g., $(4, 0)$ vs $(4, 1)$ for price hunters), but the pattern holds: discount-heavy for price hunters, quality-heavy for premium shoppers, balanced for bulk buyers.

1.5.4 0.4.4 Learning Curves and Baselines

We visualize learning progress and compare to baselines.

```

import matplotlib.pyplot as plt

def plot_learning_curves(history: List[float], window: int = 50):
    """Plot smoothed learning curve with baselines."""
    # Compute rolling average
    smoothed = np.convolve(history, np.ones(window)/window, mode='valid')

    fig, ax = plt.subplots(figsize=(10, 6))

    # Learning curve
    ax.plot(smoothed, label='Q-learning (smoothed)', linewidth=2)

    # Baselines
    random_baseline = np.mean([reward(x, A[rng.integers(len(A))])
                               for _ in range(1000)])

```

Learning Curves

Figure 1: Learning Curves

```
        for x in X])
ax.axhline(random_baseline, color='red', linestyle='--',
           label=f'Random policy ({random_baseline:.2f})')

# Static best (tuned for average user)
static_best = np.mean([reward(x, (2, 2)) for _ in range(300) for x in X])
ax.axhline(static_best, color='orange', linestyle='--',
           label=f'Static best ({static_best:.2f})')

# Oracle (knows user type, chooses optimally)
# Optimal actions: price_hunter->(4,0), premium->(0,4), bulk_buyer->(2,2)
oracle_rewards = {
    "price_hunter": np.mean([reward("price_hunter", (4, 0)) for _ in range(50)]),
    "premium": np.mean([reward("premium", (0, 4)) for _ in range(50)]),
    "bulk_buyer": np.mean([reward("bulk_buyer", (2, 2)) for _ in range(50)]),
}
oracle = np.mean(list(oracle_rewards.values()))
ax.axhline(oracle, color='green', linestyle='--',
           label=f'Oracle ({oracle:.2f})')

ax.set_xlabel('Episode')
ax.set_ylabel('Reward (smoothed)')
ax.set_title('Learning Curve: Contextual Bandit for Boost Optimization')
ax.legend()
ax.grid(alpha=0.3)

plt.tight_layout()
return fig

# Generate and save plot
fig = plot_learning_curves(hist)
fig.savefig('docs/book/ch00/learning_curves.png', dpi=150)
print("Saved learning curve to docs/book/ch00/learning_curves.png")
```

Expected output:

- **Random policy** (red dashed): ~0.0 average reward (baseline—random actions average out)
- **Static best** (orange dashed): ~0.3 (one-size-fits-all (2,2) helps bulk buyers but hurts price hunters and premium)
- **Q-learning** (blue solid): Starts near 0, converges to ~1.6 by episode 1500
- **Oracle** (green dashed): ~2.0 (theoretical maximum with perfect knowledge of optimal actions per user)

Key insight: Q-learning reaches about 82% of oracle performance by learning from experience alone. No manual tuning and no A/B tests are required. In this run, the bulk buyer segment recovers the optimal action (2, 2).

1.6 0.5 Reading the Experiment: What We Learned

1.6.1 Convergence Pattern

The learning curve has three phases:

1. **Pure exploration** (episodes 0–500): High variance, ε -greedy tries random actions, Q-values are noisy
2. **Exploitation begins** (episodes 500–1500): Agent identifies good actions per context, reward climbs steadily
3. **Convergence** (episodes 1500–3000): Q-values stabilize, reward plateaus at ~82% of oracle

This is **regret minimization** in action. Chapter 1 formalizes this; Chapter 6 analyzes convergence rates.

1.6.2 Per-Segment Performance

If we track rewards separately by user type:

```
# Track per-segment performance
segment_rewards = {x: [] for x in X}

for _ in range(100): # 100 test episodes
    for x in X:
        a = max(A, key=lambda a: Q[x][a]) # Greedy policy (no exploration)
        r = reward(x, a)
        segment_rewards[x].append(r)

for x in X:
    print(f'{x}: mean reward = {np.mean(segment_rewards[x]):.3f}')
```

Output:

```
price_hunter : mean reward = 2.309
premium      : mean reward = 2.163
bulk_buyer   : mean reward = 0.917
```

Analysis:

- **Price hunters** get the highest rewards (~2.3)—the agent found a near-optimal action (4, 1) with high discount boost
- **Premium shoppers** get high rewards (~2.2)—high quality boost (1, 4) closely matches their preferences
- **Bulk buyers** get lower rewards (~0.9) because their **balanced preferences** have inherently lower optimal reward (base=1.0 at (2,2)) compared to polarized users (base=2.5). But the agent finds the **exact optimal!**
- All three segments dramatically beat the static baseline (~0.3 average) through personalization

This is **personalization** at work: different users get different rankings, each optimized for their revealed preferences.

1.6.3 What We (Hand-Wavily) Assumed

This toy experiment “just worked,” but we made implicit assumptions:

1. **Rewards are well-defined expectations** over stochastic outcomes (Chapter 2 makes this measure-theoretically rigorous)
2. **Exploration is safe** (in production, bad rankings lose users; Chapter 9 introduces off-policy evaluation for safer testing)
3. **The logging policy and new policy have sufficient overlap** to compare fairly (importance weights finite; Chapter 9)

4. **ε -greedy tabular Q converges** (for bandits, this follows from stochastic approximation theory; Chapter 3's Bellman contraction analysis applies to the full MDP case with $\gamma > 0$)
5. **Actions are discrete and state space is tiny** (Chapter 7 handles continuous actions; Chapter 4 builds realistic state)

None of these are free. The rest of the book makes them precise and shows when they hold (or how to proceed when they don't).

1.6.4 Theory-Practice Gap: ε -Greedy Exploration

Our toy used ε -greedy exploration with constant $\varepsilon = 0.1$. This deserves scrutiny.

What theory says: In a stochastic K -armed bandit, a constant exploration rate ε forces perpetual uniform exploration and yields linear regret. If $\varepsilon_t \rightarrow 0$ with a suitable schedule, ε -greedy can achieve sublinear regret, but its exploration remains uniform over non-greedy arms. By contrast, UCB-type algorithms direct exploration through confidence bounds and achieve logarithmic (gap-dependent) regret and worst-case $\tilde{O}(\sqrt{KT})$ regret (Auer et al. 2002; Lattimore and Szepesvári 2020).

What practice shows: ε -greedy with constant $\varepsilon \in [0.05, 0.2]$ is often competitive because:

1. **Trivial to implement:** No confidence bounds, no posterior sampling, just a random number generator
2. **Handles non-stationarity gracefully:** Continues exploring even after “convergence” (useful when user preferences drift)
3. **The regret difference matters only at scale:** For the short horizons in this chapter, the gap between ε -greedy and UCB is typically negligible

When ε -greedy fails: High-dimensional action spaces where uniform exploration wastes samples. For our 25-action toy problem, it is adequate. For Chapter 7's continuous actions (10^{100} effective arms), we need structured exploration (UCB, Thompson Sampling).

Modern context: Google's 2010 display ads paper (Li et al. 2010) used ε -greedy successfully at scale. In many contemporary bandit systems, Thompson Sampling is a strong default due to its uncertainty-driven exploration and empirical performance (Russo et al. 2018; Lattimore and Szepesvári 2020).

Why UCB and Thompson Sampling? (Preview for Chapter 6)

ε -greedy explores **uniformly**—it wastes samples on arms it already knows are bad. UCB explores **optimistically**—it tries arms whose rewards *might* be high given uncertainty:

- **UCB:** Choose $a_t = \arg \max_a [Q(x, a) + \beta\sigma(x, a)]$ where σ is a confidence width. Explores arms with high uncertainty, not randomly.
- **Thompson Sampling:** Maintain posterior $P(Q^* \mid \text{data})$, sample $\tilde{Q} \sim P$, act greedily on sample. Naturally balances exploration (high posterior variance \rightarrow diverse samples) with exploitation.

Both achieve $\tilde{O}(d\sqrt{T})$ regret for d -dimensional linear bandits—matching the lower bound up to logarithms (Chu et al. 2011; Lattimore and Szepesvári 2020). In this structured setting, naive uniform exploration can be provably suboptimal, and the gap widens in high dimensions.

1.7 0.6 Limitations: Why We Need the Rest of the Book

Our toy is **pedagogical**, not production-ready. Here's what breaks at scale:

1.7.1 1. Discrete Action Space

We used 25 templates. Real search has continuous boosts: $\mathbf{w} \in [-5, 5]^{10}$ (ten features, unbounded). Discretizing to a grid would require $100^{10} = 10^{20}$ actions—tractable.

Solution: Chapter 7 introduces **continuous action bandits** via $Q(x, a)$ regression and cross-entropy method (CEM) optimization.

1.7.2 2. Tabular State Representation

We had 3 user types. Real search has thousands of user segments (RFM bins, geographic regions, device types, time-of-day). Plus query features (length, specificity, category). A realistic context space is **high-dimensional and continuous**.

Solution: Chapter 6 (neural linear bandits), Chapter 7 (deep Q-networks with continuous state/action).

1.7.3 3. No Constraints

Our agent optimized R_{toy} without guardrails. Real systems must enforce: - Profitability floors (CM2 \geq threshold) - Exposure targets (strategic products get visibility) - Rank stability (limit reordering volatility)

Solution: Chapter 10 introduces production **guardrails** (CM2 floors, $\Delta\text{Rank}@k$ stability), with Chapter 3 (Section 3.5) providing the formal CMDP theory and Lagrangian methods.

1.7.4 4. Simplified Position Bias

We didn't model how clicks depend on rank. Real users exhibit **position bias** (top-3 slots get 80% of clicks) and **abandonment** (quit after 5 results if nothing relevant).

Solution: Chapter 2 develops PBM/DBN click models; Chapter 5 implements them in `zoosim`.

1.7.5 5. Online Exploration Risk

We trained by interacting with users directly (episodes = real searches). In production, bad rankings **cost real money and lose real users**. We need safer evaluation.

Solution: Chapter 9 introduces **off-policy evaluation (OPE)**: estimate new policy performance using logged data from old policy, without deploying.

1.7.6 6. Single-Episode Horizon

We treated each search as independent. Real users return across sessions. Today's ranking affects tomorrow's retention.

Solution: Chapter 11 extends to **multi-episode MDPs** with inter-session dynamics (retention, satisfaction state).

1.8 0.7 Map to the Book

Here's how our toy connects to the rigorous treatment ahead:

Toy Concept	Formal Treatment	Chapter
User types	Context space \mathcal{X} , distribution ρ	1
Boost templates	Action space \mathcal{A} , policy π	1, 6, 7
R_{toy}	Reward function $R : \mathcal{X} \times \mathcal{A} \times \Omega \rightarrow \mathbb{R}$, constraints	1
ε -greedy Q-learning	Bellman operator, contraction mappings	3
Stochastic outcomes	Probability spaces, click models (PBM/DBN)	2
Learning curves	Regret bounds, sample complexity	6
Static best vs oracle	Importance sampling, off-policy evaluation	9
Guardrails (missing)	CMDP (Section 3.5), production guardrails	3, 10
Engagement proxy	Multi-episode MDP, retention modeling	11

Chapters 1–3 provide foundations: contextual bandits, measure theory, Bellman operators. Chapters 4–8 build the simulator and core algorithms. Chapters 9–11 handle evaluation, robustness, and production deployment. Chapters 12–15 cover frontier methods (slate ranking, offline RL, multi-objective optimization).

1.9 0.8 How to Use This Book

1.9.1 For Practitioners

We recommend working through Chapter 0 in full: we run the code, modify the reward function (Exercise 0.1), and compare exploration strategies (Exercise 0.2).

We then skim Chapters 1–3 on a first read. We focus on: - The reward formulation (#EQ-1.2 in Chapter 1) - Why engagement matters (Section 1.2.1) - The Bellman contraction intuition (Chapter 3, skip proof details initially)

We then dive into Chapters 4–11 (simulator, algorithms, evaluation). This provides the implementation roadmap.

We return to theory as needed. When something fails (e.g., divergence in a Q-network), we revisit Chapter 3’s convergence analysis.

1.9.2 For Researchers / Mathematically Inclined

We skim Chapter 0 to see the concrete thread.

We start at Chapter 1. We work through definitions, theorems, and proofs, and we verify that the code validates the mathematics.

We do the exercises: a mix of proofs (30%), implementations (40%), experiments (20%), and conceptual questions (10%).

We use Chapter 0 as a touchstone. When abstractions feel heavy, we return to the toy: “How does this theorem explain why the tabular method stabilized in Section 0.4?”

1.9.3 For Everyone

Ethos: mathematics and code are inseparable. Every theorem compiles. Every algorithm is proven rigorous, then implemented in production-quality code. Theory and practice in constant dialogue.

If a proof appears without code or code appears without theory, something is missing.

1.10 Exercises (Chapter 0)

Exercise 0.1 (Reward Sensitivity) [15 minutes]

Modify `reward()` to use different weights in R_{toy} : - (a) Pure GMV: (1.0, 0.0, 0.0) (no profitability or engagement terms) - (b) Profit-focused: (0.4, 0.5, 0.1) (prioritize CM2 over GMV) - (c) Engagement-heavy: (0.5, 0.2, 0.3) (high click weight)

For each, train Q-learning and report: - Final average reward - Learned actions per user type - Does the policy change? Why?

Hint: Case (c) risks “clickbait” strategies (see Chapter 1, Section 1.2.1). Monitor conversion quality.

Exercise 0.2 (Action Geometry) [30 minutes]

Compare two exploration strategies:

Strategy A (current): ε -greedy with uniform random action sampling

Strategy B (neighborhood): ε -greedy with **local perturbation**: when exploring, sample action near current best $a^* = \arg \max_a Q(x, a)$:

```
def explore_local(x, sigma=1.0):
    a_star = max(A, key=lambda a: Q[x][a])
    i_star, j_star = a_star
    i_new = np.clip(i_star + rng.integers(-1, 2), 0, 4)
    j_new = np.clip(j_star + rng.integers(-1, 2), 0, 4)
    return (i_new, j_new)
```

Implement both, train for 1000 episodes, and plot learning curves. Which converges faster? Why?

Reflection: This is **structured exploration**. Chapter 6 introduces UCB and Thompson Sampling, which balance exploration and exploitation more principled than ε -greedy.

Exercise 0.3 (Regret Shape) [45 minutes, extended]

Define **cumulative regret** as the gap between oracle and agent:

$$\text{Regret}(T) = \sum_{t=1}^T (R_t^* - R_t)$$

where R_t^* is the oracle reward (best action for context x_t) and R_t is the agent's reward.

(a) Implement regret tracking:

```
def compute_regret(history, contexts, oracle_Q):
    regret = []
    cumulative = 0.0
    for t, (x, r) in enumerate(zip(contexts, history)):
        r_star = oracle_Q[x]
        cumulative += (r_star - r)
        regret.append(cumulative)
    return regret
```

(b) Plot cumulative regret vs episode count. Is it sublinear (i.e., does $\text{Regret}(T)/T \rightarrow 0$)?

(c) Fit a curve: $\text{Regret}(T) \approx C\sqrt{T}$. Does this match theory? (Chapter 6 derives $O(\sqrt{T})$ regret for UCB.)

Exercise 0.4 (Advanced: Constraints) [60 minutes, extended]

Add a simple CM2 floor constraint: reject actions that violate profitability.

Setup: Modify `reward()` to return `(r, cm2)`. Define a floor $\tau = 0.3$ (30% margin minimum).

Constrained Q-learning:

```
def choose_action_constrained(x, eps, tau_cm2):
    # Filter feasible actions
    feasible = [a for a in A if expected_cm2(x, a) >= tau_cm2]
    if not feasible:
        return A[rng.integers(len(A))]  # Fallback to unconstrained

    if rng.random() < eps:
        return feasible[rng.integers(len(feasible))]

    return max(feasible, key=lambda a: Q[x][a])
```

- (a) Implement `expected_cm2(x, a)` (running average like Q).
- (b) Train with $\tau = 0.3$. How does performance change vs unconstrained?
- (c) Plot the Pareto frontier: GMV vs CM2 as τ varies over [0.0, 0.5].

Connection: This is a **Constrained MDP (CMDP)**. Chapter 3 (Section 3.5) develops the Lagrangian theory, and Chapter 10 implements production guardrails for multi-constraint optimization.

Exercise 0.5 (Bandit-Bellman Bridge) [20 minutes, conceptual]

Our toy is a **contextual bandit**: single-step decisions, no sequential states.

The **Bellman equation** (Chapter 3) for an MDP is:

$$V^*(s) = \max_a \left\{ R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^*(s') \right\}$$

where $\gamma \in [0, 1]$ is a discount factor.

Question: Show that our Q-learning update is the $\gamma = 0$ special case of Bellman.

Hint: - Set $\gamma = 0$ in Bellman equation - Note that with no future states, $V^*(s) = \max_a R(s, a)$ - Our Q-table is $Q(x, a) \approx \mathbb{E}[R | x, a]$, so $V^*(x) = \max_a Q(x, a)$ - This is **one-step value iteration**

Reflection: Contextual bandits are MDPs with horizon 1. Multi-episode search (Chapter 11) requires the full Bellman machinery.

1.11 0.9 Code Artifacts

All code from this chapter is available in the repository:

Code-Artifact Mapping

- **Run script:** `scripts/ch00/toy_problem_solution.py:1` (use `--chapter0` to reproduce this chapter's output) - **Sanity tests:** `tests/ch00/test_toy_example.py:1` (deterministic regression for Chapter 0 output) - **Learning curve plot:** `docs/book/ch00/learning_curves.png:1` (generated artifact)

To reproduce:

```
uv run python scripts/ch00/toy_problem_solution.py --chapter0
uv run pytest -q tests/ch00
```

Expected output:

```
Final average reward (last 100 episodes): 1.640
```

Learned policy:

<code>price_hunter</code>	\rightarrow action (4, 1) (Q = 1.948)
<code>premium</code>	\rightarrow action (1, 4) (Q = 2.289)
<code>bulk_buyer</code>	\rightarrow action (2, 2) (Q = 0.942)

```
Saved learning curve to docs/book/ch00/learning_curves.png
```

1.12 0.10 What's Next?

We have now trained a first RL agent for search ranking. It learned context-adaptive boost weights from scratch, achieving near-oracle performance without manual tuning.

But we cheated. We used a tiny discrete action space, three user types, and online exploration without safety guarantees. Real systems need:

1. **Rigorous foundations** (Chapters 1–3): Formalize contextual bandits, measure-theoretic probability, Bellman operators
2. **Realistic simulation** (Chapters 4–5): Scalable catalog generation, position bias models, rich user dynamics
3. **Continuous actions** (Chapter 7): Regression-based Q-learning, CEM optimization, trust regions
4. **Constraints and guardrails** (Chapter 10): CM2 floors, Δ Rank@k stability, safe fallback policies
5. **Safe evaluation** (Chapter 9): Off-policy evaluation (IPS, DR, FQE) for production deployment
6. **Multi-episode dynamics** (Chapter 11): Retention modeling, long-term value, engagement as state

The journey from toy to production is the journey of this book.

In **Chapter 1**, we formalize everything we hand-waved here: What exactly is a contextual bandit? Why is the reward function EQ-1.2 mathematically sound? How do constraints become a CMDP? Why does engagement matter, and when should it be implicit vs explicit?

Let us make it rigorous.

End of Chapter 0

Auer, Peter, Nicolò Cesa-Bianchi, and Paul Fischer. 2002. “Finite-Time Analysis of the Multiarmed Bandit Problem.” *Machine Learning* 47 (2–3): 235–56.

Chu, Wei, Lihong Li, Lev Reyzin, and Robert E. Schapire. 2011. “Contextual Bandits with Linear Payoff Functions.” *Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics (AISTATS)*, JMLR workshop and conference proceedings, vol. 15: 208–14.

Lattimore, Tor, and Csaba Szepesvári. 2020. *Bandit Algorithms*. Cambridge University Press. <https://doi.org/10.1017/9781108571401>.

Li, Lihong, Wei Chu, John Langford, and Robert E. Schapire. 2010. “A Contextual-Bandit Approach to Personalized News Article Recommendation.” *Proceedings of the 19th International Conference on World Wide Web (WWW)*, 661–70.

Robbins, Herbert, and Sutton Monro. 1951. “A Stochastic Approximation Method.” *The Annals of Mathematical Statistics* 22 (3): 400–407.

Russo, Daniel J., Benjamin Van Roy, Abbas Kazerouni, Ian Osband, and Zheng Wen. 2018. “A Tutorial on Thompson Sampling.” *Foundations and Trends in Machine Learning* 11 (1): 1–96. <https://doi.org/10.1561/2200000070>.